

EMBSYS 105

Programming with Embedded & Real-Time Operating Systems

Instructor: Nick Strathy, nstrathy@uw.edu

TA: Gideon Lee, gideonhlee@yahoo.com

© N. Strathy 2020

Lecture 4


1/27/2020

Looking ahead

Date	Lecture number	Assignment
1/6	L1	A1 due* before L2
1/13	L2	A2 due before L3
1/20	L3	A3 due before L4
1/27	L4	A4 due before L5
2/3	L5	A5 due before L7, Project due before L10
2/10	Holiday (?)	
2/17	L6	
2/24	L7	
3/2	L8	
3/9	L9	
3/16	L10 – Student presentations	

* Assignments are due Sunday night at 11:59 PM

Previous Lecture (L3) Overview

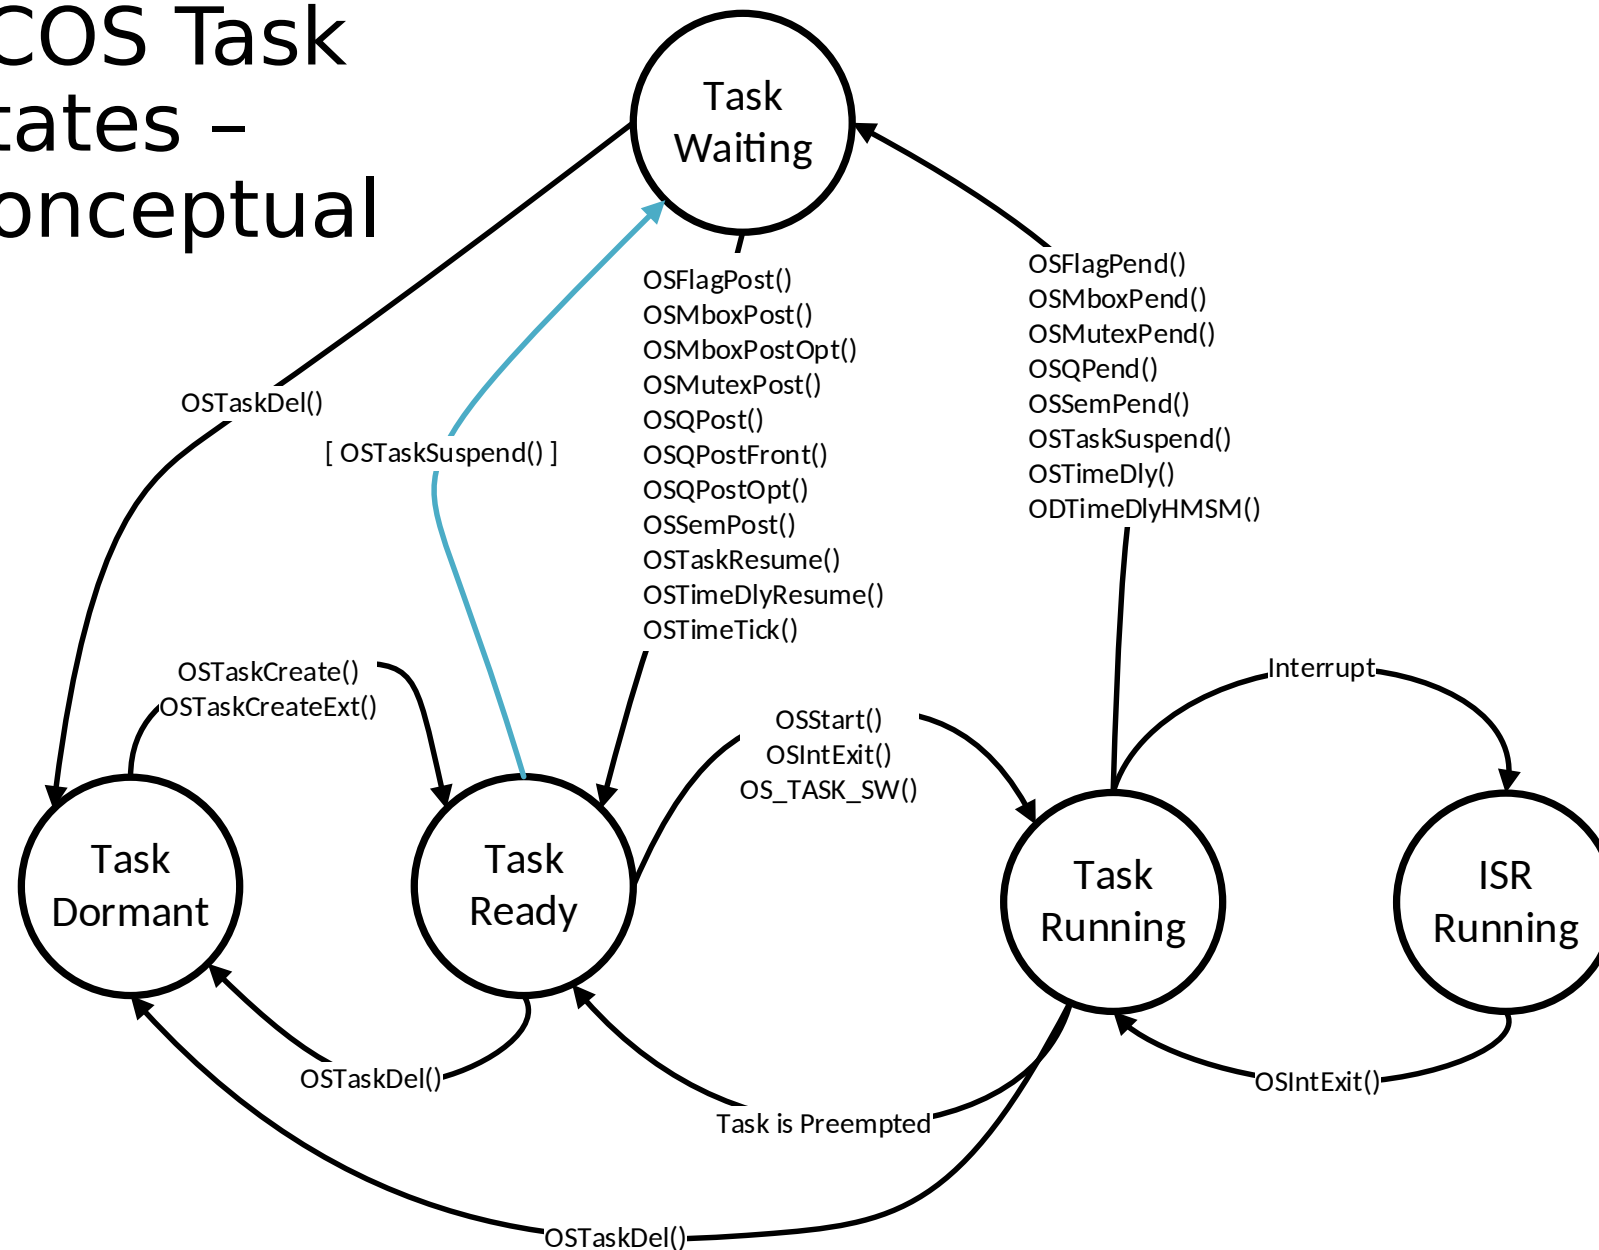
- MicroC/OS-II (uCOS) Introduction
 - Task creation, Task Delay, Sample task code, uCOS startup steps
 - Porting uCOS to our board (**Labrosse Ch 13**)
 - Key data definitions, enabling and disabling interrupts, critical section implementation, initializing the stack, context switching, C pointers and assembly language
 - uCOS Services (**Labrosse Ch 16**)
 - Task Management
 - Time Management
 - Interrupt Management
 - Semaphores
 - Mutexes
 - Message Mailboxes
 - Message Queues
 - Event Flags
 - Memory Management
 - User-Defined Functions
 - Miscellaneous Services
 - uCOS Configuration (**Labrosse Ch 17**)
 - Assignment 3 – uCOS Port
- 
- Stopped here last lecture.
Next slide: 64

Current Lecture (L4) Overview

First, finish slides from last week (L3) beginning at slide 64

- uCOS Internals (Labrosse ch 3, 6, 9)
 - Task States
 - Task Control Blocks (TCBs)
 - Ready List
 - Task Scheduling
 - Event Control Blocks
 - Placing a task in the ECB Wait List
 - Removing a Task from the ECB Wait List
 - Find the Highest Priority Task Waiting on the ECB
 - List of Free ECBs
 - Initializing an ECB
 - Making a Task Ready
 - Making a Task Wait for an Event
 - Making a task Ready Because of a Timeout
 - Event Flags
- Task Synchronization Techniques (Tanenbaum, Wikipedia, etc.)
 - Sharing data between ISRs and tasks
 - Producer-Consumer Problem
 - Readers and Writers Problem
 - Deadlock
 - Dining Philosophers Problem
- Event Driven Systems (Wikipedia)
 - Characteristics of Event Driven Systems
 - Data Flow Diagrams
- Assignment 4 – Task Synchronization

uCOS Task States - Conceptual



Notes:

- Labrosse p. 39 does not show Ready \rightarrow Waiting however `OSTaskSuspend(prio)` can make that transition.
- These states are conceptual – actually there are several Wait states; no actual Running or Dormant states are *tracked* in later versions of uCOS.
- Later versions of uCOS call “ISR Running” “Interrupted”

uCOS Task States – Implementation

```
INT8U          OSTCBStat;          /* Task      status          */
```

The above TCB field takes on values chosen from:

```
#define OS_STAT_RDY          0x00u    /* Ready to run          */
#define OS_STAT_SEM          0x01u    /* Pending on semaphore  */
#define OS_STAT_MBOX         0x02u    /* Pending on mailbox     */
#define OS_STAT_Q            0x04u    /* Pending on queue       */
#define OS_STAT_SUSPEND      0x08u    /* Task is suspended     */
#define OS_STAT_MUTEX        0x10u    /* Pending on mutual exclusion semaphore */
#define OS_STAT_FLAG         0x20u    /* Pending on event flag group */

#define OS_STAT_PEND_ANY      (OS_STAT_SEM | OS_STAT_MBOX | OS_STAT_Q | OS_STAT_MUTEX | OS_STAT_FLAG)
```

```
INT8U          OSTCBStatPend;      /* Task PEND status      */
```

The above TCB field takes on values chosen from:

```
#define OS_STAT_PEND_OK      0u        /* Pending status OK, not pending, or pending complete */
#define OS_STAT_PEND_TO      1u        /* Pending timed out */
#define OS_STAT_PEND_ABORT    2u        /* Pending aborted */
```

Task Control Block (OS_TCB)

```
typedef struct os_tcb {
    OS_STK          *OSTCBStkPtr;          /* Pointer to current top of stack */
                                                */
#ifdef OS_TASK_CREATE_EXT_EN > 0
    void            *OSTCBExtPtr;          /* Pointer to user definable data for TCB extension */
    OS_STK          *OSTCBStkBottom;       /* Pointer to bottom of stack */
    INT32U          OSTCBStkSize;          /* Size of task stack (in number of stack elements) */
    INT16U          OSTCBOpt;              /* Task options as passed by OSTaskCreateExt() */
    INT16U          OSTCBId;               /* Task ID (0..65535) */
#endif

    struct os_tcb   *OSTCBNext;            /* Pointer to next TCB in the TCB list */
    struct os_tcb   *OSTCBPrev;            /* Pointer to previous TCB in the TCB list */

#ifdef OS_EVENT_EN || (OS_FLAG_EN > 0)
    OS_EVENT         *OSTCBEventPtr;       /* Pointer to event control block */
#endif

#ifdef ((OS_Q_EN > 0) && (OS_MAX_QS > 0)) || (OS_MBOX_EN > 0)
    void            *OSTCBMsg;             /* Message received from OSMboxPost() or OSQPost() */
#endif
#endif
```

Task Control Block (OS_TCB)

```
#if (OS_FLAG_EN > 0) && (OS_MAX_FLAGS > 0)
#if OS_TASK_DEL_EN > 0
    OS_FLAG_NODE    *OSTCBFlagNode;    /* Pointer to event flag node */
#endif
    OS_FLAGS        OSTCBFlagsRdy;    /* Event flags that made task ready to run */
#endif

    INT16U          OSTCBDly;          /* Nbr ticks to delay task or, timeout waiting for event */
    INT8U           OSTCBStat;          /* Task status */
    INT8U           OSTCBStatPend;      /* Task PEND status */
    INT8U           OSTCBPrio;          /* Task priority (0 == highest) */

    INT8U           OSTCBX;             /* Bit position in group corresponding to task priority */
    INT8U           OSTCBY;             /* Index into ready table corresponding to task priority */
#if OS_LOWEST_PRIO <= 63
    INT8U           OSTCBBitX;          /* Bit mask to access bit position in ready table */
    INT8U           OSTCBBitY;          /* Bit mask to access bit position in ready group */
#else
    INT16U          OSTCBBitX;          /* Bit mask to access bit position in ready table */
    INT16U          OSTCBBitY;          /* Bit mask to access bit position in ready group */
#endif
#endif
```


Task Control Block (OS_TCB)

```
#if OS_TASK_DEL_EN > 0
    INT8U          OSTCBDelReq;      /* Indicates whether a task needs to delete itself */
#endif

#if OS_TASK_PROFILE_EN > 0
    INT32U          OSTCBCtxSwCtr;    /* Number of time the task was switched in */
    INT32U          OSTCBCyclesTot;   /* Total number of clock cycles the task has been running */
    INT32U          OSTCBCyclesStart; /* Snapshot of cycle counter at start of task resumption */
    OS_STK          *OSTCBStkBase;   /* Pointer to the beginning of the task stack */
    INT32U          OSTCBStkUsed;    /* Number of bytes used from the stack */
#endif

#if OS_TASK_NAME_SIZE > 1
    INT8U          OSTCBTaskName[OS_TASK_NAME_SIZE];
#endif
} OS_TCB;
```

TCB Free List (ucos-ii.h)

```
OS_EXT  OS_TCB      *OSTCBCur;           /* Pointer to currently running TCB */
OS_EXT  OS_TCB      *OSTCBFreeList;      /* Pointer to list of free TCBs */
OS_EXT  OS_TCB      *OSTCBHighRdy;       /* Pointer to highest priority TCB R-to-R */
OS_EXT  OS_TCB      *OSTCBLList;         /* Pointer to doubly linked list of TCBs */
OS_EXT  OS_TCB      *OSTCBPrioTbl[OS_LOWEST_PRIO + 1]; /* Table of pointers to created TCBs */
OS_EXT  OS_TCB      OSTCBTbl[OS_MAX_TASKS + OS_N_SYS_TASKS]; /* Table of TCBs */
```

- OSTCBFreeList
 - New tasks get their TCB from this list
 - The new TCB moves from the free list to OSTCBLList
- OSTCBPrioTbl
 - New tasks get their TCB added to this table indexed by priority
 - If there is already a non-null entry for a given priority, a task with that priority already exists and error code OS_ERR_PRIO_EXIST is returned by OSTaskCreate()
- All TCBs are preallocated at compile time in array OSTCBTbl

Ready List

```
#define OS_LOWEST_PRIO 31
```

```
#define OS_RDY_TBL_SIZE ((OS_LOWEST_PRIO) / 8 + 1)
```

```
OS_EXT INT8U OSRdyTbl[OS_RDY_TBL_SIZE];
```

- Rather than maintain a sorted list of Ready tasks with the highest priority task always at the front, uCOS represents the Ready list by a set of lookup tables based on the task priority
- This scheme results in constant time for Ready List insertion/deletion regardless of the number of tasks. It also yields constant time for determining the highest priority Ready task regardless of the number of tasks
- The tables used to maintain the Ready list are **OSRdyTbl** , **OSRdyGrp**, and **OSUnMapTbl**
- **OSRdyTbl** contains the set of tasks that are in the Ready state **indexed by task priority**
- Each task is represented by 1 bit in OSRdyTbl where 0 means not Ready, 1 means Ready

Ready List

Add a task to the Ready List (table):

```
OSRdyTbl[ptcb->OSTCBY] |= ptcb->OSTCBBitX;
```

Where these are computed once during task creation:

```
ptcb->OSTCBY      = (INT8U)(prio >> 3);  
ptcb->OSTCBX      = (INT8U)(prio & 0x07);  
  
ptcb->OSTCBBitY   = (INT8U)(1 << ptcb->OSTCBY);  
ptcb->OSTCBBitX   = (INT8U)(1 << ptcb->OSTCBX);
```

Remove a task from the Ready List (table):

```
y      = ptcb->OSTCBY;  
OSRdyTbl[y]      &= ~ptcb->OSTCBBitX;
```

Ex: Make Task 20 ready:

```
prio 2010 == 101002 binary  
ptcb->OSTCBY = 102 = 210  
ptcb->OSTCBX = 1002 = 410  
ptcb->OSTCBBitX = 100002  
ptcb->OSRdyTbl[2] |= 100002
```

OSRdyTbl:

<div>x y</div>	7	6	5	4	3	2	1	0
0	0	0	0	0	1	0	0	0
1	0	0	0	0	0	1	0	0
2	0	0	0	1	0	0	0	0
3	0	0	1	0	1	0	0	0

Ready List

Finding the highest priority Ready Task

```
INT8U          OSRdyGrp; /* global, tracks high-order 3 bits of Ready task priorities */  
  
OSRdyGrp      |= ptcb->OSTCBBitY; /* updated whenever a task becomes Ready/unReady */
```

Note: OSRdyGrp has bit i set if row i of OSRdyTbl has any bits set.

Where these are computed once during task creation (see previous slide):

```
ptcb->OSTCBY      = (INT8U)(prio >> 3);  
ptcb->OSTCBBitY   = (INT8U)(1 << ptcb->OSTCBY);
```

Use OSRdyGrp, OSRdyTbl, and OSUnMapTbl to get highest priority Ready task:

```
INT8U y          = OSUnMapTbl[OSRdyGrp];  
OSPrioHighRdy = (INT8U)((y << 3) + OSUnMapTbl[OSRdyTbl[y]]);
```

Ready List

OSUnMapTable enables lookup of highest priority Ready task given OSRdyGrp and OSRdyTbl

```
INT8U y          = OSUnMapTbl[OSRdyGrp];
OSPrioHighRdy = (INT8U)((y << 3) + OSUnMapTbl[OSRdyTbl[y]]);
where:
INT8U  const  OSUnMapTbl[256] = {
    0, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x00 to 0x0F */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x10 to 0x1F */
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x20 to 0x2F */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x30 to 0x3F */
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x40 to 0x4F */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x50 to 0x5F */
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x60 to 0x6F */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x70 to 0x7F */
    7, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x80 to 0x8F */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0x90 to 0x9F */
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0xA0 to 0xAF */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0xB0 to 0xBF */
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0xC0 to 0xCF */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0xD0 to 0xDF */
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, /* 0xE0 to 0xEF */
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0 /* 0xF0 to 0xFF */
};
```

Task Scheduling

- OS_Sched() performs task-level scheduling
 - Determines the highest priority ready task
 - Transfers control of the CPU to that task
 - Note: uCOS doesn't explicitly change state from Ready \rightarrow Running i.e. the running task remains in the Ready List.
- OSIntExit() performs ISR-level scheduling
 - Only reschedules to the highest priority ready task if interrupt nesting level is 0 otherwise simply returns from procedure call to nested ISR code which will return from ISR to previous nested ISR

OSSched()

```
void OS_Sched (void)
{
    #if OS_CRITICAL_METHOD == 3                /* Allocate storage for CPU status register */
        OS_CPU_SR cpu_sr = 0;
    #endif

    OS_ENTER_CRITICAL();
    if (OSIntNesting == 0) {                  /* Schedule only if all ISRs done and ... */
        if (OSLockNesting == 0) {            /* ... scheduler is not locked */
            OS_SchedNew(); /* OS_SchedNew() determines highest priority Ready task */
            if (OSPrioHighRdy != OSPrioCur) { /* No Ctx Sw if current task is highest rdy */
                OSTCBHighRdy = OSTCBPrioTbl[OSPrioHighRdy];
            #if OS_TASK_PROFILE_EN > 0
                OSTCBHighRdy->OSTCBCtxSwCtr++; /* Inc. # of context switches to this task */
            #endif
            OSCtxSwCtr++;                      /* Increment context switch counter */
            OS_TASK_SW(); /* Perform a context switch - really a call to OSCtxSw() */
        }
    }
    OS_EXIT_CRITICAL();
}
```


OSSchedNew() – finds highest priority Ready task

```
static void OS_SchedNew (void)
{
    #if OS_LOWEST_PRIO <= 63                /* See if we support up to 64 tasks */
        INT8U    y;

        y          = OSUnMapTbl[OSRdyGrp];
        OSPrioHighRdy = (INT8U)((y < 3) + OSUnMapTbl[OSRdyTbl[y]]);
    #else                                  /* We support up to 256 tasks */
        INT8U    y;
        INT16U *ptbl;

        if ((OSRdyGrp & 0xFF) != 0) {
            y = OSUnMapTbl[OSRdyGrp & 0xFF];
        } else {
            y = OSUnMapTbl[(OSRdyGrp >> 8) & 0xFF] + 8;
        }
        ptbl = &OSRdyTbl[y];
        if ((*ptbl & 0xFF) != 0) {
            OSPrioHighRdy = (INT8U)((y < 4) + OSUnMapTbl[(*ptbl & 0xFF)]);
        } else {
            OSPrioHighRdy = (INT8U)((y < 4) + OSUnMapTbl[(*ptbl >> 8) & 0xFF] + 8);
        }
    #endif
}
```

OSIntExit()

```
void OSIntExit (void)
{
    #if OS_CRITICAL_METHOD == 3                /* Allocate storage for CPU status register */
        OS_CPU_SR cpu_sr = 0;
    #endif
    if (OSRunning == OS_TRUE) {
        OS_ENTER_CRITICAL();
        if (OSIntNesting > 0) {                /* Prevent OSIntNesting from wrapping */
            OSIntNesting--;
        }
        if (OSIntNesting == 0) {                /* Reschedule only if all ISRs complete ... */
            if (OSLockNesting == 0) {           /* ... and not locked. */
                OS_SchedNew();
                if (OSPrioHighRdy != OSPrioCur) { /* No Ctx Sw if current task is highest rdy */
                    OSTCBHighRdy = OSTCBPrioTbl[OSPrioHighRdy];
                }
            }
            #if OS_TASK_PROFILE_EN > 0
                OSTCBHighRdy->OSTCBCtxSwCtr++; /* Inc. # of context switches to this task */
            #endif
            OSCtxSwCtr++;                        /* Keep track of the number of ctx switches */
            OSIntCtxSw();                        /* Perform interrupt level ctx switch */
        }
    }
    OS_EXIT_CRITICAL();
}
```

Event Control Blocks (OS_EVENTS)

- ECBs are used for managing 4 types of uCOS task synchronization “subclasses”
 - Semaphores
 - Mutexes
 - Message Mailboxes
 - Message Queues
- We looked already at the APIs for operating on ECBs
- Now we'll look at the internal data and operations
- Data consists of
 - ECB type (“subclass”)
 - Data specific to the ECB type
 - Wait list for tasks blocked on the ECB

Event Control Blocks (OS_EVENTS)

```
typedef struct os_event {  
    INT8U    OSEventType;          /* Type of event control block (see OS_EVENT_TYPE_xxxx) */  
    void     *OSEventPtr;          /* Pointer to message or queue structure */  
    INT16U    OSEventCnt;          /* Semaphore Count (not used if other EVENT type) */  
#if OS_LOWEST_PRIO <= 63  
    INT8U     OSEventGrp;          /* Group corresponding to tasks waiting for event to occur */  
    INT8U     OSEventTbl[OS_EVENT_TBL_SIZE]; /* List of tasks waiting for event to occur */  
#else  
    INT16U     OSEventGrp;          /* Group corresponding to tasks waiting for event to occur */  
    INT16U     OSEventTbl[OS_EVENT_TBL_SIZE]; /* List of tasks waiting for event to occur */  
#endif  
  
#if OS_EVENT_NAME_SIZE > 1  
    INT8U     OSEventName[OS_EVENT_NAME_SIZE];  
#endif  
} OS_EVENT;
```

Event Control Blocks (OS_EVENTS)

- OSEventType (“subclass”) is one of
 - OS_EVENT_TYPE_SEM, OS_EVENT_TYPE_MUTEX, OS_EVENT_TYPE_MBOX, OS_EVENT_TYPE_Q
- OSEventPtr points to the message if this is a Mailbox, else a queue structure if this is a Queue
- OSEventCnt is the semaphore counter if this is a semaphore
- OSEventTbl[] and OSEventGrp together implement the wait list using the same scheme as the Ready List.
- OSEventName[] is the name of the ECB.
 - Use OSEventNameSet(), OSEventNameGet()

Event Control Blocks (OS_EVENTS)

- Adding/Removing a task from an ECB wait list
 - This needs to be done when blocking/unblocking a task on the ECB
 - Analogous to adding/removing a task from the Ready List
 - Uses OSEventGrp and OSEventTbl of the OS_EVENT instance instead of OSRdyGrp and OSRdyTbl
 - Note: each ECB has its own OSEventGrp and OSEventTbl
- Finding the highest priority task waiting on an ECB
 - This needs to be done when deciding which task to unblock
 - Analogous to finding the highest priority task in the Ready List
 - Uses OSEventGrp, OSEventTbl and again OSUnMapTbl

Event Control Blocks (OS_EVENTS)

Making a task wait on an ECB

- Below function is called by OSSemPend(), OSMboxPend(), etc. to block the calling task on the ECB

```
void OS_EventTaskWait (OS_EVENT *pevent)
{
    INT8U y;

    OSTCBCur->OSTCBEventPtr = pevent; /* Store pointer to event control block in TCB */
    y = OSTCBCur->OSTCBY; /* Task no longer ready */
    OSRdyTbl[y] &= ~OSTCBCur->OSTCBBitX;
    if (OSRdyTbl[y] == 0) {
        OSRdyGrp &= ~OSTCBCur->OSTCBBitY; /* Clear bit if this was only task ready in grp*/
    }
    pevent->OSEventTbl[OSTCBCur->OSTCBY] |= OSTCBCur->OSTCBBitX; /* Put task in wait list*/
    pevent->OSEventGrp |= OSTCBCur->OSTCBBitY;
}
```

Event Control Blocks (OS_EVENTS)

Making a task ready because of a Timeout

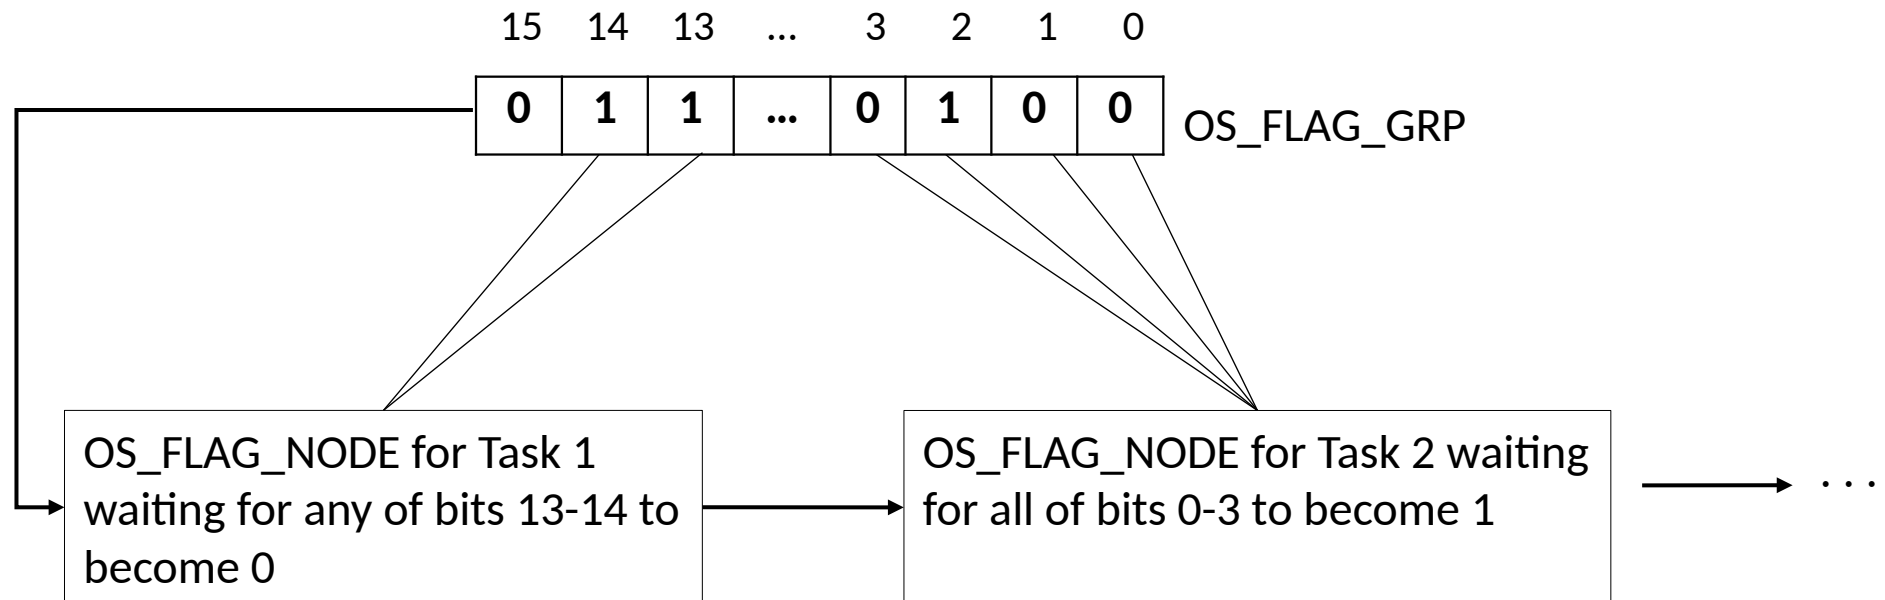
- When we do a blocking wait **with timeout**, ptcb->OSTCBDly is loaded with the timeout ticks
- OSTimeTick() decrements OSTCBDly and if it reaches 0, makes the task Ready
- Below function is called by OSSemPend(), OSMboxPend(), etc. to handle timeout on blocking wait

```
void OS_EventT0Abort (OS_EVENT *pevent)
{
    INT8U y;
    y = OSTCBCur->OSTCBY;
    pevent->OSEventTbl[y] &= ~OSTCBCur->OSTCBBitX; /* Remove task from wait list */
    if (pevent->OSEventTbl[y] == 0x00) {
        pevent->OSEventGrp &= ~OSTCBCur->OSTCBBitY;
    }
    OSTCBCur->OSTCBStatPend = OS_STAT_PEND_OK; /* Clear pend status */
    OSTCBCur->OSTCBStat = OS_STAT_RDY; /* Set status to ready */
    OSTCBCur->OSTCBEventPtr = (OS_EVENT *)0; /* No longer waiting for event */
}
```


Event Flags

High level view of Event Flag implementation (example)

- Implemented as an OS_FLAG_GRP struct which points to a list of OS_FLAG_NODE structs
- Each OS_FLAG_NODE represents a task waiting for a combination of bits in the OS_FLAG_GRP



Event Flags

OS_FLAG_GRP typedef

```
typedef struct os_flag_grp {          /* Event Flag Group */
    INT8U      OSFlagType;           /* Should be set to OS_EVENT_TYPE_FLAG */
    void      *OSFlagWaitList;       /* Pointer to first NODE of task waiting on event flag */
    OS_FLAGS    OSFlagFlags;         /* 8, 16 or 32 bit flags */
#ifdef OS_FLAG_NAME_SIZE > 1
    INT8U      OSFlagName[OS_FLAG_NAME_SIZE];
#endif
} OS_FLAG_GRP;
```

Notes:

- OSFlagType must be OS_EVENT_TYPE_FLAG (i.e. not OS_EVENT_TYPE_SEM, OS_EVENT_TYPE_MBOX, etc.)
- OSFlagWaitList points to a doubly linked list of OS_FLAG_NODES (didn't find a good reason why it's declared void*)
- OS_FlagFlags is the 16-bit set of flags in this flag group

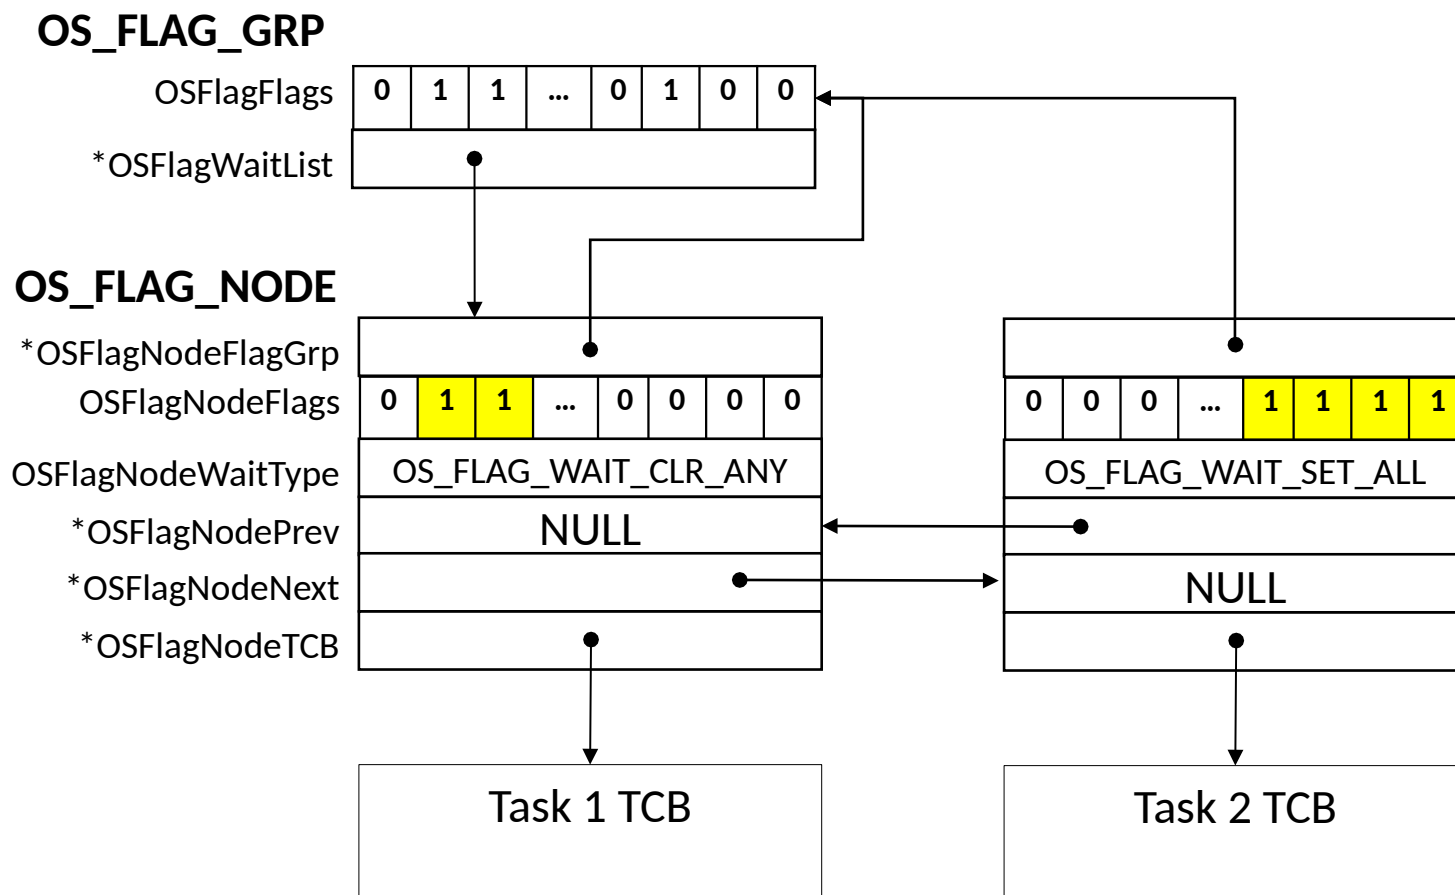
Flag Events

OS_FLAG_NODE typedef

```
typedef struct os_flag_node {  
    void        *OSFlagNodeNext;  
    void        *OSFlagNodePrev;  
    void        *OSFlagNodeTCB;  
    void        *OSFlagNodeFlagGrp;  
    OS_FLAGS    OSFlagNodeFlags;  
    INT8U       OSFlagNodeWaitType;  
  
    /* Event Flag Wait List Node */  
    /* Pointer to next      NODE in wait list */  
    /* Pointer to previous NODE in wait list */  
    /* Pointer to TCB of waiting task */  
    /* Pointer to Event Flag Group */  
    /* Event flag to wait on */  
    /* Type of wait: */  
    /*      OS_FLAG_WAIT_AND */  
    /*      OS_FLAG_WAIT_ALL */  
    /*      OS_FLAG_WAIT_OR */  
    /*      OS_FLAG_WAIT_ANY */  
  
} OS_FLAG_NODE;
```

Note: Didn't find a good reason why pointers are all void*

Flag Events



Waiting on a Flag Group

- Allocate an OS_FLAG_NODE from free list and insert it at front of OSFlagWaitList
- Initialize the fields and point OSFlagNodeTCB to the current TCB
- Remove the current TCB from Ready List and change its status to OS_STAT_FLAG
- Reschedule

Posting to a Flag Group

- Update the OSFlagFlags in OS_FLAG_GRP with the posted value
- Chain through OSFlagWaitList and for any OS_FLAG_NODES whose bit patterns are satisfied:
 - Make Ready the linked TCB
 - Delete the OS_FLAG_NODE and return it to free list
- Reschedule

Summary of uCOS Internals

- Task States

- uCOS keeps the Running task in the Ready List
- Tracks several types of Wait (blocked) states
- Does not track Dormant, Running, or Interrupted states

- Task Control Blocks (TCBs)

- The TCB contains all the information needed to manage the Task
- Are all preallocated at compile time
- Are initialized individually at Task Creation time

- Ready List

- Implemented as a set of tables which allow insertion, deletion, and lookup of the highest priority task in constant time regardless of number of tasks.

- Task Scheduling

- Handled by two routines:
- OSSched() handles context switch from one task to another
- OSIntExit() handles context switch from ISR code to the highest priority task but only if interrupt nesting level is 0

- Event Control Blocks

- Provide a uniform API and management framework for 4 “subclasses” of ECB: semaphores, mutexes, message mailboxes, and message queues.
- Repurpose the Ready List table lookup scheme to track tasks blocked on any “subclass” instance

- Event Flags

- Enable tasks to block for arbitrary events signaled by other tasks or ISRs
- Implemented as a Flag Group node of reference bits referenced by a doubly linked list of blocked TCBs waiting for a TCB-specific bit pattern to occur in the Flag Group.

Task Sync Techniques

- We will look at some common issues for synchronizing tasks
- The goal is to have a bag of tricks for handling synchronization problems
- Synchronization techniques must
 - Avoid lengthy ISRs
 - Avoid shared resource corruption when multiple tasks operate on a shared resource simultaneously
 - Avoid deadlock – when two or more tasks are blocked waiting for a resource held by one of the other tasks.
 - Avoid starvation – where one or more tasks are starved for access to a resource

Task Sync Techniques

Issues around sharing resources between ISRs and tasks

- Typical scenario:
 - interrupt occurs when new data is ready to be processed
 - ISR moves the data to a shared location (buffer) for a task to process
 - How to synchronize the ISR and task to guard the buffer?
- ISRs cannot do blocking waits for resources (uCOS prevents that by returning an error code if you try).
- ISRs should obviously not spin while waiting for a resource
- ISRs simply cannot wait. Period.
- So what should they do?

Task Sync Techniques

A wrong way: ISR moves new data to the buffer and posts to a semaphore to have a task process the new data

- Task code blocks and waits for a post to allow it to process new data
- ISR code moves new data to a shared buffer, then posts to the semaphore to allow it to process the input
- Note: ISRs can't use uCOS Mutexes: only task code can acquire a mutex and only the owner can release it

ISR Code

```
ISR() {  
    // Move new data to the  
    // shared buffer  
    Post(mySem);  
}
```

Task Code

```
while (1) {  
    Pend(mySem);  
    // Remove data from the  
    // shared buffer;  
    // process the data  
}
```


Task Sync Techniques

Problem:

- We're using a semaphore to protect our shared resource, so we should be good, right? Wrong! Semaphores always require careful thought
- What happens if another interrupt from the same source occurs while the task is accessing the buffer?
- In general we are not going to keep interrupts disabled while the task is processing the previous data
- How can we guard the shared buffer?

ISR Code

```
ISR() {  
    // Move new data to the  
    // shared buffer  
    Post(mySem);  
}
```

Task Code

```
while (1) {  
    Pend(mySem);  
    // Remove data from the  
    // shared buffer;  
    // process the data  
}
```

Task Sync Techniques

A Solution: Use a message passing scheme

- Use a message passing scheme such as a mailbox to pass the buffer pointer back and forth between task and ISR.
- Now if an interrupt from the same source happens before the task has removed the previous message, we'll lose the new data but at least we won't corrupt the old data
- If we can't afford to lose the new data, use a sufficiently long **message queue** instead of a mailbox.

ISR Code

```
ISR() {  
    if (Accept(mBox, msg)) {  
        // move new data to the  
        // msg buffer  
        Post(mBox, msg);  
    }  
}
```

Task Code

```
while (1) {  
    msg = Pend(mBox);  
    // process the msg  
    Post(mBox, msg);  
}
```

Task Sync Techniques

Producer-Consumer Problem

- In the absence of message passing services, we can use semaphores
- The classical Producer-Consumer Problem involves synchronizing a Producer Task which deposits data in a buffer and a Consumer Task which takes data out of the same buffer
- With no synchronization, the result is chaos with the producer overwriting the previous buffer contents while the consumer is removing data

Producer Task Code

```
while (1) {  
    // get new data  
    // deposit data in buffer  
}
```

Consumer Task Code

```
While (1) {  
    // remove data from buffer  
    // process data  
}
```

Task Sync Techniques

Producer-Consumer Problem

- Can we protect the buffer with just one semaphore?
- Uh-oh, that's no good – while the Producer is getting new data the Consumer can try to remove a message multiple times
- While the Consumer is processing data the Producer can overwrite the buffer with new data multiple times

Producer Task Code

```
while (1) {  
    // get new data  
    Pend(guardSem);  
    // deposit data in  
    // buffer  
    Post(guardSem);  
}
```

Consumer Task Code

```
While (1) {  
    Pend(guardSem);  
    // remove data from  
    // buffer  
    Post(guardSem);  
    // process data  
}
```

Task Sync Techniques

Producer-Consumer Problem

- How about adding another semaphore to block the Producer till the Consumer is ready for new data?
- Initialization
 - guardSem <- 1
 - emptySem <- 1
- Uh-oh the Consumer can still spin removing the same message multiple times while the Producer is getting new data

Producer Task Code

```
while (1) {  
    // get new data  
    Pend(emptySem);  
    Pend(guardSem);  
    // deposit data in  
    // buffer  
    Post(guardSem);  
}
```

Consumer Task Code

```
While (1) {  
    Pend(guardSem);  
    // remove data from  
    // buffer  
    Post(guardSem);  
    // process data  
    Post(emptySem);  
}
```

Task Sync Techniques

Producer-Consumer Problem

- How about one more semaphore to block the consumer till the Producer has new data?
- Initialization
 - guardSem <- 1
 - emptySem <- 1
 - fullSem <- 0
- Now we're OK.

Producer Task Code

```
while (1) {  
    // get new data  
    Pend(emptySem);  
    Pend(guardSem);  
    // deposit data in  
    // buffer  
    Post(guardSem);  
    Post(fullSem);  
}
```

Consumer Task Code

```
While (1) {  
    Pend(fullSem);  
    Pend(guardSem);  
    // remove data from  
    // buffer  
    Post(guardSem);  
    Post(emptySem);  
    // process data  
}
```

Task Sync Techniques

Producer-Consumer Problem

- Can generalize this to where the message buffer is a queue with **N** slots
- Initialization
 - `guardSem <- 1`
 - `fullSem <- 0`
 - **`emptySem <- N`**
- Allows up to N messages to be deposited in the queue before blocking the Producer until the Consumer removes a message
- Meanwhile the Consumer can remove a message any time `fullSem` is > 0

Producer Task Code

```
while (1) {  
    // get new data  
    Pend(emptySem);  
    Pend(guardSem);  
    // deposit data in  
    // queue  
    Post(guardSem);  
    Post(fullSem);  
}
```

Consumer Task Code

```
While (1) {  
    Pend(fullSem);  
    Pend(guardSem);  
    // remove data from  
    // queue  
    Post(guardSem);  
    Post(emptySem);  
    // process data  
}
```

Task Sync Techniques

Producer-Consumer Problem

- How about if the Producer is an ISR?
- ISRs can't do blocking waits. Can we do it with non-blocking waits?
- **No: can't use this scheme as-is if Producer is an ISR**
- There may be available slots in the queue but if the Consumer is actively removing data from the queue the guardSem semaphore will not be available and the ISR will have to throw away the new data
- **Solution:** rather than use guardSem, disable interrupts while adding/removing data from the queue ...

ISR Producer Code

```
ISR() {  
    // get new data  
    if (Accept(emptySem)) {  
        if (Accept(guardSem)) {  
            // deposit data in  
            // queue  
            Post(guardSem);  
            Post(fullSem);  
        }  
        else Post(emptySem);  
    }  
}
```

Task Consumer Code

```
While (1) {  
    Pend(fullSem);  
    Pend(guardSem);  
    // remove data from  
    // queue  
    Post(guardSem);  
    Post(emptySem);  
    // process data  
}
```


Task Sync Techniques

Producer-Consumer Problem

- Solution if Producer is an ISR
- Deposit and removal from the queue must be done with interrupts off
- Otherwise if we use a semaphore the consumer can be active in the critical section thus preventing the ISR from operating on the queue

ISR Producer Code

```
ISR() {  
    // get new data  
    if (Accept(emptySem)) {  
        DisableInterrupts();  
        // deposit data in  
        // queue  
        EnableInterrupts();  
        Post(fullSem);  
    }  
}
```

Task Consumer Code

```
While (1) {  
    Pend(fullSem);  
    DisableInterrupts();  
    // remove data from  
    // queue  
    EnableInterrupts();  
    Post(emptySem);  
    // process data  
}
```

Task Sync Techniques

Readers and Writers Problem

- This problem has multiple readers reading from shared data who need to synchronize with writers who update the shared data
- If no writer is active, any number of readers may read the data simultaneously
- Only one writer may update the data at a time and no readers may read while the writer is writing
- Initialize
 - `rc <- 0` (int reader count)
- This solution (from Tanenbaum) may starve writers. True for uCOS?
- Other solutions can be found in the literature

Writers Code

```
while (1) {  
    // get new data  
    Pend(wMutex);  
    // write new data  
    // to shared data  
    Post(wMutex);  
}
```

Readers Code

```
While (1) {  
    Pend(rcMutex);  
    rc += 1;  
    if (rc == 1) Pend(wMutex);  
    Post(rcMutex);  
    // read shared data  
    Pend(rcMutex);  
    rc -= 1;  
    if (rc == 0) Post(wMutex);  
    Post(rcMutex);  
    // process the data  
}
```

Task Sync Techniques

Deadlock

- Deadlock or Deadly Embrace can happen when multiple tasks try to acquire a set of resources in different orders:
 - Task 1 gets mutexA
 - Task1 gets preempted by Task 2
 - Task2 gets mutexB
 - deadlock
- One way to reduce the chance of deadlock is to always acquire in name-sorted order
- Easier said than done if one semaphore is acquired at one place in the code and a second one later at some distant place in the code.

Task 1

```
while (1) {  
    Pend(mutexA);  
    Pend(mutexB);  
    // use resources A and B  
    Post(mutexB);  
    Post(mutexA);  
}
```

Task 2


```
while (1) {  
    Pend(mutexB);  
    Pend(mutexA);  
    // use resources A and B  
    Post(mutexA);  
    Post(mutexB);  
}
```

Task Sync Techniques

Non solution:

Deadlock/Starvation Avoidance

- **Dining Philosophers Problem**
- N philosophers each spend their time in an endless loop: think; eat; repeat;
- The catch is that chopsticks are shared resources



```
void philosopher(int i) {  
    while (1) {  
        think();  
        takeChopstick(i);  
        takeChopstick((i+1) % N);  
        eat();  
        putChopStick(i);  
        putChopStick((i+1) % N);  
    }  
}
```

- Solutions to such problems can require much thought to avoid both deadlock in contention for resources and starvation where one task rarely or never gets access to shared resources.
- There are numerous solutions to the Dining Philosophers. See Tanenbaum, Wikipedia, etc.

Task Sync Techniques

Summary of Task Sync Techniques

- Wrong ways and right ways to share data between ISRs and tasks
- Producer-Consumer Problem
 - Task-based exploration
 - ISR-as-Producer exploration
- Readers and Writers Problem
- Deadlock and Starvation
 - Example deadlock due to incorrect acquisition order of mutexes
 - Name-sorted-order acquisition of mutexes as deadlock avoidance technique
 - Dining Philosophers Problem

Event Driven Systems

An Event Driven System is composed of

- Event Emitters (or Agents)
 - responsible for detecting, gathering and transferring events
- Event Consumers (or Sinks)
 - responsible for providing a reaction to a received event. The reaction may consist of filtering and forwarding the filtered event
- Event Channels
 - conduct the events from Emitter to Consumer

The Graphical User Interface is the canonical example of an Event Driven System

- Event Channel is a message passing queue where Emitters post messages to the queue and Consumers consume messages from the queue
- Emitters are keyboard, mouse
 - Events are key-press messages, mouse movement- and click-messages
- Consumers are windows which block for messages placed in the event channel by Emitters

Event Driven Systems

Data Flow Diagrams (DFDs)

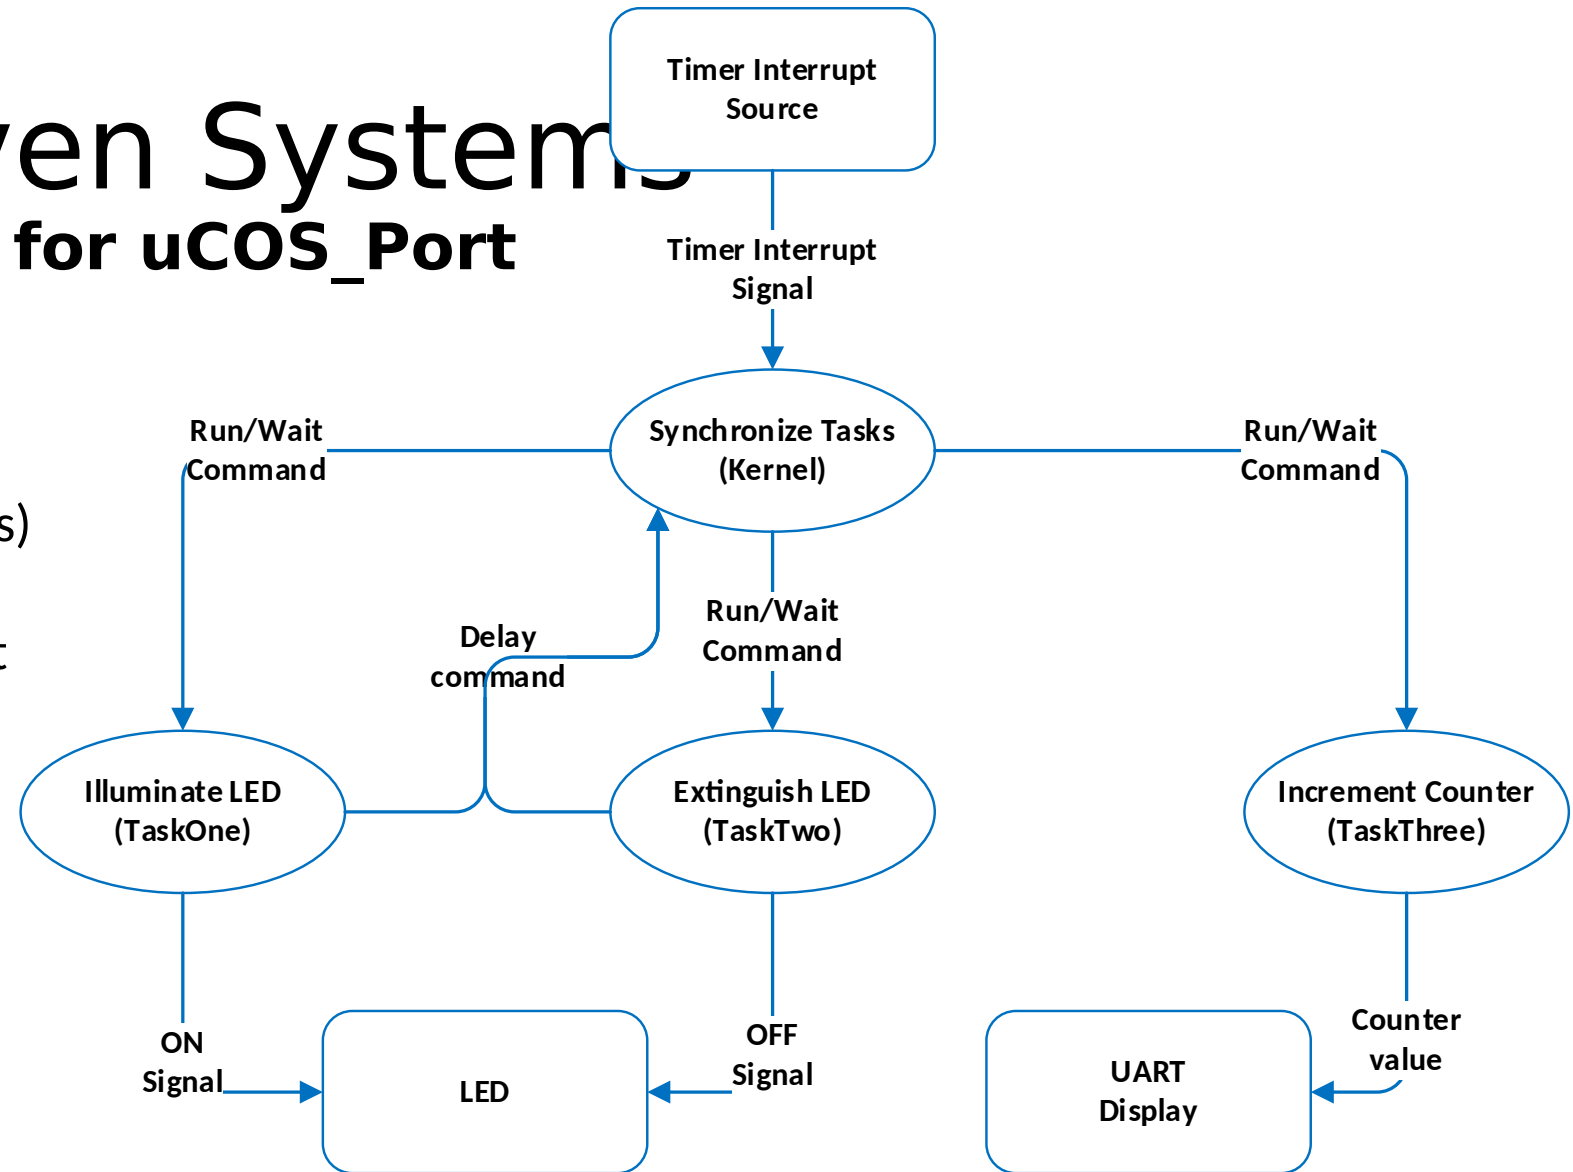


- A useful abstraction tool for many kinds of software systems (not just Event Driven Systems)
- Composed of 4 basic primitives: Terminators, Data stores, Processes, and Data Flows
- Less detailed than Flow Charts
- Useful at the design level for thinking through the design
- Not a specification but a way of explaining the behavior of the system
- In Embedded Event Driven Systems, *Process* nodes can often map to Tasks and *Data Flow* links can map to passed messages/command-codes

Event Driven Systems

DFD Abstraction for uCOS_Port

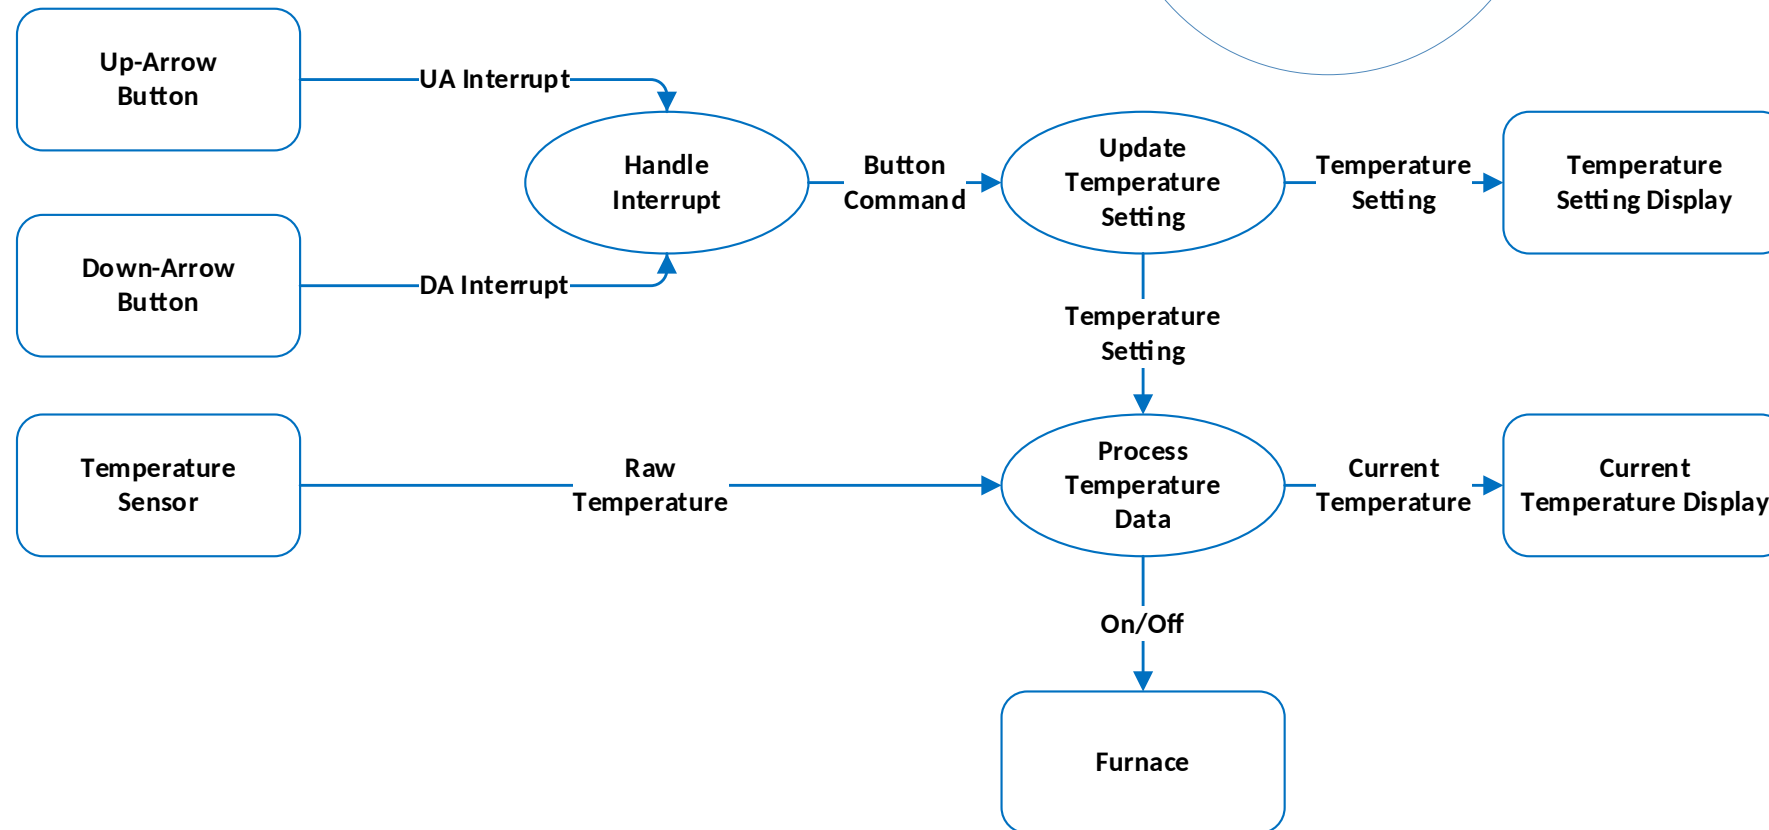

- Processes (ovals) get imperative verb labels
- Terminators (rectangles) get noun labels
- Data flows (arrows) get noun labels



Event Driven System

DFD for button-driven system

- interrupt driven buttons



Event Driven Systems

DFD for button-driven system

- No button interrupts (polling of b

